

COMP 421: Lab # 2

The Yalnix Operating System Kernel

Important Dates:

- In class, Thursday, February 16, 2012: Project begins.
- 11:59PM, Friday, March 9, 2012: Recommended project checkpoint (see Section 4.4).
- 11:59PM, Tuesday, March 20, 2012: Project completion deadline.

1. Project Overview

Through this assignment, you will be able to learn how a real operating system kernel works, managing the hardware resources of the computer system and providing services to user processes running on the system. In the project, you will implement an operating system kernel for the Yalnix operating system, running on a mythical computer system known as the Rice Computer Systems RCS 421. The Yalnix operating system is similar in many ways to the Unix operating system, with influences also from the V-System, a research operating system written at Stanford University. In particular, Yalnix supports multiple processes, each having their own virtual address space, plus interprocess communication between the processes based on message passing. The system also supports a system console and a set of terminals on which users can do input and output.

Through the magic of the support software provided for your use in this project, a single Unix process will be made to behave like an RCS 421 computer for you to run your Yalnix kernel on. That is, when you run your kernel, it will appear that you are running your own operating system on a real RCS 421 computer, but in reality, everything will be running in user mode in a single Unix process. As in the first project, separate X terminal windows will also be used to simulate the terminals attached to the RCS 421 computer.

The project will be done in groups of 2 students. You are free to choose your own partner for the project. The third project in the class will also be done in groups of 2 students; you may keep the same partner for the third project or change partners between these two projects if you like.

In working on the project with your partner, make sure that you each know what the other is doing, and think ahead to design data structures and approaches that can be shared by both of you. In order to get the full benefit of the course, both partners in a project group should fully understand your group's solution to the project.

The project must be done on the CLEAR Linux system at Rice. In particular, the hardware simulation of the RCS 421 computer system requires certain libraries and executable files that are only installed on CLEAR and are only compatible with CLEAR. You can access the CLEAR system by doing an `ssh` to the machine name `ssh.clear.rice.edu` and logging in using your NetID as your password.

2. The RCS 421 Computer System

This section describes the hardware architecture of the Rice Computer Systems RCS 421 computer, a mythical computer architecture on which you will implement an operating system kernel in this project. The operating system, known as Yalnix, is described in Section 3. The RCS 421 computer system has the features of real hardware but is simplified somewhat to make the project more manageable.

This section is written generally in the style of a real computer system hardware architecture manual. Everything in this section is something that a real operating system designer and implementor on such a computer would need to know about the hardware. Certain features of the operating system running on top

of this hardware are mentioned here only as guidance in the implementation of any operating system for this hardware.

2.1. Machine Registers

This section defines the general purpose registers and privileged registers of the RCS 421 computer system.

2.1.1. General Purpose Registers

The RCS 421 computer has a number of general purpose registers that may be accessed from either user mode or kernel mode. These general purpose registers include the following:

- PC, the *program counter*, contains the virtual address from which the currently executing instruction was fetched.
- SP, the *stack pointer*, contains the virtual address of the top of the current process's stack.
- R0 through R7, the eight *general registers*, used as accumulators or otherwise to hold temporary values during a computation.

In reality, the real hardware has quite a few additional general purpose registers, but since you will be writing your operating system in C or C++, not in machine language or assembly language, only these registers are relevant to the project.

2.1.2. Privileged Registers

In addition to the general purpose CPU registers, the RCS 421 computer contains a number of privileged registers, accessible only from kernel mode. Table 1 summarizes these privileged registers.

Details for the use of each of these registers are described in the following sections, as indicated in Table 1, where the operation of that component of the RCS 421 computer system hardware architecture is defined. Most of the privileged registers are both writable and readable, except for the REG_TLB_FLUSH register, which is write-only.

All values written into or read from these privileged registers must be of type RCS421RegVal. This data type is defined in the file comp421/hardware.h. The comp421/hardware.h include file is located in the directory /clear/courses/comp421/pub/include/comp421. You can include the file comp421/hardware.h in your source file by using

```
#include <comp421/hardware.h>
```

The hardware provides two privileged instructions for reading and writing these privileged machine registers:

Table 1 RCS 421 Privileged Register Summary

Name	Purpose	Readable	Details
REG_PTR0	Page Table Register 0	Yes	Section 2.2.4
REG_PTR1	Page Table Register 1	Yes	Section 2.2.4
REG_TLB_FLUSH	TLB Flush Register	No	Section 2.2.5
REG_VM_ENABLE	Virtual Memory Enable Register	Yes	Section 2.2.6
REG_VECTOR_BASE	Vector Base Register	Yes	Section 2.4.1

- `void WriteRegister(int which, RCS421RegVal value)`
Write value into the privileged machine register designated by `which`.
- `RCS421RegVal ReadRegister(int which)`
Read the register specified by `which` and return its current value.

Each privileged machine register is identified by a unique integer constant as noted in Table 1, passed as the `which` argument to the instructions. The file `comp421/hardware.h` defines the symbolic names for these constants.

As noted above, the values of these registers are represented by the data type `RCS421RegVal`. You must use a C or C++ “cast” to convert other data types (such as addresses) to type `RCS421RegVal` when calling `WriteRegister`, and must also use a “cast” to convert the value returned by `ReadRegister` to the desired type if you need to interpret the returned value as any type other than an `RCS421RegVal`.

2.2. Memory Subsystem

2.2.1. Overview

The memory subsystem of the RCS 421 computer supports physical memory size that depends on the amount of hardware memory installed. The physical memory is divided into pages of size `PAGESIZE` bytes.

The individual pages of this physical memory may be mapped into the address space of a running process through virtual memory supported by the hardware and initialized and controlled by your operating system. The hardware component that implements the hardware control for these virtual memory mappings is known as the *Memory Management Unit*, or *MMU*. As with the physical memory, the virtual memory is divided into pages of size `PAGESIZE`, and any page of the virtual address space may be mapped to any page of physical memory through the *page table* of each running process.

The constant `PAGESIZE` is defined in `comp421/hardware.h`. In addition, this include file also defines a number of other constants and macros to make dealing with addresses and page numbers easier:

- `PAGESIZE`: The size in bytes of each physical memory and virtual memory page.
- `PAGEOFFSET`: A bit mask that can be used to extract the byte offset with a page, given an arbitrary physical or virtual address. For an address `addr`, the value `(addr & PAGEOFFSET)` is the byte offset within the page where `addr` points.
- `PAGEMASK`: A bit mask that can be used to extract the beginning address of a page, given an arbitrary physical or virtual address. For an address `addr`, the value `(addr & PAGEMASK)` is the beginning address of the page where `addr` points.
- `PAGESHIFT`: The log base 2 of `PAGESIZE`, which is thus also the number of bits in the offset in a physical or virtual address.
- `UP_TO_PAGE(addr)`: A C preprocessor macro that rounds address `addr` up to the next highest page boundary. If `addr` is already on a page boundary, it returns `addr`.
- `DOWN_TO_PAGE(addr)`: A C preprocessor macro that rounds address `addr` down to the next lowest page boundary. If `addr` is already on a page boundary, it returns `addr`.

Due to hardware restrictions (in reality, due to restrictions in the underlying Unix operating system), the bottom `MEM_INVALID_SIZE` bytes (`MEM_INVALID_PAGES` pages) of the address space are inaccessible. Before virtual memory has been enabled, this means that any physical address in this range of addresses will cause an error. After virtual memory has been enabled, this means that all page table entries (`MEM_INVALID_PAGES` page table entries) corresponding to this address range must be all have a `valid` bit of 0. By enforcing this restriction, the hardware attempts to ensure that accidental accesses to memory using a `NULL` pointer are detectable.

2.2.2. Physical Memory

The beginning address and size of physical memory in the RCS 421 computer system are defined as follows:

- `PMEM_BASE`: The physical memory address of the first byte of physical memory in the machine. This address is a constant and is determined by the machine's hardware design. The value `PMEM_BASE` is defined in `comp421/hardware.h`.
- The total size (number of bytes) of physical memory in the computer is determined by how much RAM is installed on the machine. At boot time, the computer's firmware tests the physical memory and determines the amount of memory installed, and passes this value to the initialization procedure of the operating system kernel. The size of physical memory is unrelated to the size of the virtual address space of the machine.

As described in Section 2.2.1, the physical memory of the machine is divided into pages of size `PAGESIZE`. Page numbers in physical memory addresses are often referred to as *page frame numbers*.

2.2.3. Virtual Address Space

The virtual address space of the machine is divided into two regions, called *Region 0* and *Region 1*. Region 0 is used by user processes executing on the operating system, and Region 1 is used by the operating system kernel. In other words, Region 0 will contain the state of each process, and Region 1 will be managed by your kernel to hold kernel state.

This division of the virtual address space into two regions is illustrated in Figure 1. The beginning virtual address of Region 0 is defined as `VMEM_0_BASE`, the limit virtual address of Region 0 is defined as `VMEM_0_LIMIT`, and the size of Region 0 is defined as `VMEM_0_SIZE`. Likewise, the beginning virtual address of Region 1 is defined as `VMEM_1_BASE`, the limit virtual address is `VMEM_1_LIMIT`, and the size is `VMEM_1_SIZE`. By the definition of the hardware, the two regions are adjacent, such that `VMEM_1_BASE` equals `VMEM_0_LIMIT`. The overall beginning virtual address of virtual memory is `VMEM_BASE` (which equals `VMEM_0_BASE`), and the overall limit virtual address of virtual memory is `VMEM_LIMIT` (which equals `VMEM_1_LIMIT`).

The state of the kernel in Region 1 consists of two parts: the kernel's text, and the kernel's global variables (its data, bss, and heap). This state is the same, independent of which user-level process is executing. The kernel stack however, holds the local kernel state associated with the user-level process that is currently executing. Both the kernel stack and all of Region 0, are switched on a context switch to map the context of the processes being switched. The hardware provides the two regions so that it can protect kernel data structures from illegal tampering by the user applications: when the CPU is executing in kernel mode, references to addresses in either region are allowed, but when the CPU is executing in user mode, references to addresses in Region 1 are not allowed. The hardware enforces this restriction.

These regions are not "segments" (e.g., as defined in Section 8.6 of the course textbook). Actual segments have an associated base address and size that are managed (and changeable) by the operating

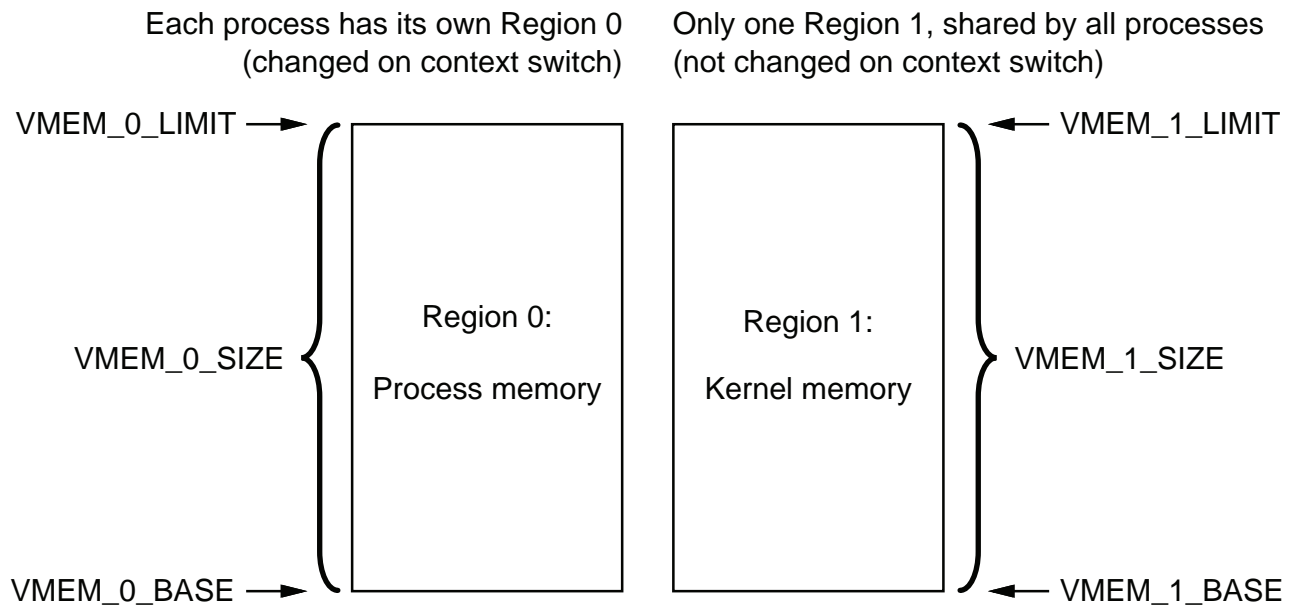


Figure 1 RCS 421 Hardware Virtual Memory Address Region Layout

system, whereas the regions in the hardware in this project have locations and sizes that are fixed in the virtual address space by the MMU hardware.

2.2.4. Page Tables

The MMU hardware automatically translates each reference (use) of any virtual memory address into the corresponding physical memory addresses. This translation happens automatically in the hardware, without intervention from the kernel. The kernel does, however, control the mapping that specifies this translation, by building and initializing the contents of the page table. This division of labor makes sense because memory references occur very frequently as the CPU executes instructions, whereas the virtual-to-physical address mapping (the page table) changes relatively infrequently and in accord with the process abstraction and memory protection policy implemented by the kernel.

The hardware uses a direct, single-level page table to define the mapping from virtual to physical addresses. Such a page table is an array of page table entries, laid out *contiguously* in physical memory. Each page table entry (PTE) contains information relevant to the mapping of a single virtual page to a corresponding physical page.

The kernel allocates page tables in memory and it tells the MMU in the hardware where to find these page tables through the following privileged registers:

- `REG_PTR0`: Contains the *physical* memory address of the beginning of the page table for Region 0 of virtual memory.
- `REG_PTR1`: Contains the *physical* memory address of the beginning of the page table for Region 1 of virtual memory.

The MMU hardware never writes to these registers. The kernel can change the virtual-to-physical address mappings by changing these registers (e.g., during a context switch) and may change any individual page table entry by modifying that entry in the corresponding page table in memory. Remember that the

value written to REG_PTR0 and REG_PTR1 must be the *physical* address of the corresponding page table. For your operating system kernel software to be able to access a page table, you will need to map it to *some* virtual address, but you must tell the hardware the *physical* address of the page table. Also, the entire page table must be *contiguous* in *physical* memory, beginning at the physical address given in this register.

When a program, or the kernel itself, makes a memory reference to any address, the MMU hardware finds the corresponding page frame number (PFN) by looking in the page table for the appropriate region of virtual memory. The page tables are indexed by virtual page number (VPN) within the region and contain page table entries (PTEs) that specify the corresponding page frame number (among other things) of each page of virtual memory in that region.

For example, suppose that PAGESIZE is defined to be 0x1000 (this is actually the value defined in `comp421/hardware.h`). If the CPU references virtual address 0x2468, this virtual address is at offset 0x468 within virtual page number 0x2 within Region 0. The MMU hardware in this case looks at the page table entry at subscript 0x2 in the page table pointed to by the physical address currently in register REG_PTR0.

Suppose also that VMEM_1_BASE is defined to be 0x200000 (this is actually the value defined in `comp421/hardware.h`). If the CPU references virtual address 0x202468, this virtual address is at offset 0x468 within virtual page number 0x2 *within* Region 1. The MMU hardware in this case looks at the page table entry at subscript 0x2 in the page table pointed to by the physical address currently in register REG_PTR1. Notice that the subscript used is the virtual page number *within the region* in which the virtual address is located, so the subscript used is 0x2, *not* 0x202.

This lookup to translate a virtual page number into its currently mapped corresponding physical page frame number is carried out automatically, wholly in hardware. In carrying out this lookup, the hardware needs to examine page table entries in order to index into the page table (the size of a page table entry must be known to the hardware) and also to extract the page frame number from the entry (the format of a page table entry must also be known to the hardware).

The format of the page table entries is dictated by the hardware. A page table entry is 32-bits wide, and contains the following defined fields:

- `valid` (1 bit): If this bit is set, the page table entry is valid; otherwise, a memory exception is generated when/if this virtual memory page is accessed. All other bits in a PTE are ignored if the `valid` bit is off.
- `kprot` (3 bits): These three bits define the memory protection applied by the hardware to this virtual memory page whenever operating in kernel mode. These three protection bits are interpreted independently as follows:
 - `PROT_READ`: Memory within the page may be read.
 - `PROT_WRITE`: Memory within the page may be written.
 - `PROT_EXEC`: Memory within the page may be executed as machine instructions.

Execution of instructions by the CPU requires that their virtual memory pages be mapped with both `PROT_READ` and `PROT_EXEC` protection set.

- `uprot` (3 bits): These three bits define the memory protection applied by the hardware to this virtual memory page whenever operating in user mode. These three protection bits are interpreted independently, in the same way as for the kernel mode protection bits (`kprot`) defined above.
- `PFN` (20 bits): This field contains the page frame number (the physical memory page number) of the page of physical memory to which this virtual memory page is mapped by this page table entry.




Figure 2 RCS 421 Hardware Page Table Entry (PTE) Format

This page table entry format is illustrated in Figure 2 and is defined in the file `comp421/hardware.h` as a C data structure as a `struct pte`. This C structure has the same memory layout as a hardware page table entry and can be used in your operating system kernel.

2.2.5. Translation Lookaside Buffer (TLB)

The RCS 421 contains a *translation lookaside buffer*, or *TLB*, to speed up virtual-to-physical address translation. The TLB caches page table entries so that subsequent references to the same virtual page do not have to retrieve the corresponding page table entry anew from physical memory.

Page table entries are loaded automatically, as needed, into the TLB by the hardware during virtual address to physical address translation. The operating system cannot directly load page table entries into the TLB and cannot examine the contents of the TLB to determine what page table entries are currently present in the TLB.

However, the operating system kernel must at times flush all or part of the TLB. In particular, after changing a page table entry in memory, the TLB may contain a stale copy of this page table entry (the value of the PTE before it was changed), since this entry may have been previously used and thus loaded into the TLB by the hardware. Also, when performing a context switch from one process to another, the operating system must flush all entries from the TLB corresponding to the virtual address space of the outgoing (old) process, since outgoing and incoming process use different page tables.

The hardware provides the `REG_TLB_FLUSH` privileged machine register to allow the operating system kernel to control the flushing of all or part of the TLB. When writing a value to the `REG_TLB_FLUSH` register, the MMU interprets the value as follows:

- `TLB_FLUSH_ALL`: Flush all entries in the entire TLB.
- `TLB_FLUSH_0`: Flush all entries in the TLB that correspond to virtual addresses in Region 0.
- `TLB_FLUSH_1`: Flush all entries in the TLB that correspond to virtual addresses in Region 1.
- Otherwise, the value written into the `REG_TLB_FLUSH` register is interpreted by the hardware as a virtual address. If an entry exists in the TLB corresponding to the virtual memory page into which this address points, flush this single entry from the TLB; if no such entry exists in the TLB, this operation has no affect.

Symbolic constants for the `TLB_FLUSH_ALL`, `TLB_FLUSH_0`, and `TLB_FLUSH_1` values are defined in `comp421/hardware.h`.

2.2.6. Enabling Virtual Memory

When the machine is booted, a boot ROM firmware loads the kernel into physical memory and begins executing it. The kernel is loaded into contiguous physical memory starting at addresses `PMEM_BASE`.

When the computer is first turned on, the machine's virtual memory subsystem is initially disabled and remains disabled until initialized and explicitly enabled by the operating system kernel. In order to initialize the virtual memory subsystem, the kernel must, for example, create an initial set of page table entries and must write the page table base and limit registers to tell the hardware where the page tables are located in memory. Until virtual memory is enabled, all addresses used are interpreted by the hardware as *physical addresses*; after virtual memory is enabled, all addresses used are interpreted by the hardware as *virtual addresses*.

The hardware provides a privileged machine register, the Virtual Memory Enable Register, that may be set by the kernel to enable the virtual memory subsystem in the hardware. To enable virtual memory, the kernel executes:

```
WriteRegister(REG_VM_ENABLE, 1)
```

After execution of this instruction, all addresses used are interpreted by the MMU hardware only as virtual memory addresses. Virtual memory cannot be disabled after it is enabled, without rebooting the machine.

Before enabling virtual memory as described above, your operating system kernel must properly initialize the page table for Region 1, which maps the kernel's own memory. The details on this process are given in Section 3.4.3.

2.3. Hardware Devices

The RCS 421 computer system can be configured with a variety of I/O and other hardware devices. For purposes of this project, only two different types of devices will be available: a number of terminals and a hardware clock.

2.3.1. Terminals

The system is equipped with `NUM_TERMINALS` terminals, numbered 0 through `NUM_TERMINALS - 1`, for use by user processes. By convention, the first terminal (terminal number 0) is considered to be the system console, while the other terminals are general-purpose terminal devices. However, this distinction is only a convention and does not imply any difference in behavior. Like in the first project, the terminals are line-oriented, but here, this behavior is implemented for you; imagine that this line-oriented behavior is implemented in the hardware (although this is not entirely realistic, it does simplify the project).

When reading from a terminal, an interrupt is generated by the hardware only after a complete line has been typed at the terminal. For writing to a terminal, a buffer of output characters (which may actually be only part of a line or may be several lines of output) is given to the hardware terminal controller, and an interrupt is not generated until that entire buffer of output has been sent to the terminal. The constant `TERMINAL_MAX_LINE`, defined in `comp421/hardware.h`, defines the maximum line length (maximum buffer size) supported for either input or output on the terminals.

In a real system, the kernel would have to manipulate the terminal device hardware registers to read and write from the device. For simplicity in this project, these details have been abstracted into two C functions:

- `void TtyTransmit(int tty_id, void *buf, int len)`

The `TtyTransmit` operation begins the transmission of `len` characters from memory, starting at *virtual* address `buf`, to terminal `tty_id`. The address `buf` must be in the kernel's virtual memory, (i.e., it must be in virtual memory Region 1), since otherwise, this memory could become unavailable

to the hardware terminal controller after a context switch. When the data has been completely written out, the terminal controller will generate a `TRAP_TTY_TRANSMIT` interrupt. When the `TRAP_TTY_TRANSMIT` interrupt handler in the operating system kernel is called, the terminal number of the terminal generating the interrupt will be made available to the handler. Until receiving a `TRAP_TTY_TRANSMIT` interrupt for this terminal after executing a `TtyTransmit` on it, no new `TtyTransmit` may be executed for this terminal.

- `int TtyReceive(int tty_id, void *buf, int len)`

When the user completes an input line on a terminal by typing either newline (`'\n'`) or return (`'\r'`), the hardware terminal controller will generate a `TRAP_TTY_RECEIVE` interrupt for this terminal (on input, both newline and return are translated to newline). The terminal number of the terminal generating the interrupt will be made available to the kernel's interrupt handler for this type of interrupt. In the interrupt handler, the kernel should execute a `TtyReceive` operation for this terminal, in order to retrieve the new input line from the hardware. The new input line is copied from the hardware for terminal `tty_id` into the kernel buffer at *virtual* address `buf`, for maximum length to copy of `len` bytes. The value of `len` *must* be equal to `TERMINAL_MAX_LINE` bytes, since the kernel cannot otherwise know the length of the new input line before calling `TtyReceive`. The buffer must be in the kernel's virtual memory (i.e., it must be entirely within virtual memory Region 1). After each `TRAP_TTY_RECEIVE` interrupt, the kernel must do a `TtyReceive` and save the new input line in a buffer inside the kernel, e.g., until a user process requests the next line from the terminal by executing a kernel call to read from this device.

The actual length of the new input line, including the newline (`'\n'`), is returned as the return value of `TtyReceive`. Thus when a blank line is typed, `TtyReceive` will return a 1, since the blank line is terminated by a newline character. When an end of file character (control-D) is typed, `TtyReceive` returns 0 for this line. End of file behaves just like any other line of input, however. In particular, you can continue to read more lines after an end of file. The data copied into your buffer by `TtyReceive` is *not* terminated with a null character (as would be typical for a string in C); to determine the end of the characters returned in the buffer, you must use the length returned by `TtyReceive`.

In reality, the `buff` address passed to `TtyTransmit` or `TtyReceive` would be interpreted by the hardware as a *physical* address. However, to simplify the project, on the RCS 421, this address is interpreted as a *virtual* address. Each terminal in the project is simulated using a separate X window, similar to the terminal behavior in the first project.

The operation of interrupts on the RCS 421 computer system is explained in more detail in Section 2.4.

2.3.2. The Hardware Clock

The RCS 421 computer system is equipped with a hardware clock that generates periodic interrupts to the operating system kernel. When a clock interrupt occurs, the hardware generates a `TRAP_CLOCK` interrupt.

The operating system must use these `TRAP_CLOCK` interrupts for any timing needs of the operating system. For example, these periodic interrupts must be used as the time source for counting down the quantum in CPU process scheduling.

The operation of interrupts on the RCS 421 computer system is explained in more detail in Section 2.4.

2.4. Interrupts, Exceptions, and Traps

2.4.1. Interrupt Vector Table and Types

Interrupts, exceptions, and traps are all handled in a similar way by the hardware, switching into kernel mode and calling a handler procedure in the kernel by indexing into the *interrupt vector table* based on the type of trap. The following types of interrupts, exceptions, and traps are defined by the RCS 421 hardware:

- `TRAP_KERNEL`: This type of trap results from a “kernel call” (also called a “system call” or “syscall”) trap instruction executed by the current user processes. Such a trap is used by user processes to request some type of service from the operating system kernel, such as creating a new process, allocating memory, or performing I/O. All different kernel call requests enter the kernel through this single type of trap.
- `TRAP_CLOCK`: This type of interrupt results from the machine’s hardware clock, which generates periodic clock interrupts.
- `TRAP_ILLEGAL`: This type of exception results from the execution of an illegal instruction by the currently executing user process. An illegal instruction can be, for example, an undefined machine language opcode or an illegal addressing mode.
- `TRAP_MEMORY`: This type of exception results from an attempt to access memory at some virtual address, when the access is disallowed by the hardware. The access may be disallowed because the address is outside the virtual address range of the hardware (outside Region 0 and Region 1), because the address is not mapped in the current page tables, or because the access violates the current page protection specified in the corresponding page table entry.
- `TRAP_MATH`: This type of exception results from any arithmetic error from an instruction executed by the current user process, such as division by zero or an arithmetic overflow.
- `TRAP_TTY_RECEIVE`: This type of interrupt is generated by the terminal device controller hardware when a line of input is completed from one of the terminals attached to the system.
- `TRAP_TTY_TRANSMIT`: This type of interrupt is generated by the terminal device controller hardware when the current buffer of data previously given to the controller on a `TtyTransmit` instruction has been completely sent to the terminal.

The interrupt vector table is stored in memory as an array of pointers to functions, each of which handles the corresponding type of interrupt, exception, or trap. The name of each type listed above is defined in `comp421/hardware.h` as a symbolic constant. This number is used by the hardware to index into the interrupt vector table to find the pointer to the handler function for this trap.

The privileged machine register `REG_VECTOR_BASE` is used by the hardware as the base *physical* address in memory where the interrupt vector table is stored. After initializing the table in memory, the kernel must write the *physical* address of the table to the `REG_VECTOR_BASE` register. The entire interrupt vector table must be *contiguous* in *physical* memory, beginning at the physical address given in the `REG_VECTOR_BASE` register.

The interrupt vector table *must* have exactly `TRAP_VECTOR_SIZE` entries in it, even though the use for some of them is currently undefined by the hardware. Any entries in the interrupt vector table other than those defined above must be initialized to `NULL`.

To simplify the programming of the operating system kernel, the kernel is not interruptible. That is, while executing inside the kernel, the kernel cannot be interrupted by any interrupt. Thus, the kernel need not use any special synchronization procedures such as semaphores or monitors inside the kernel. Any

interrupts that occur while already executing inside the kernel are held pending by the hardware and will be triggered only once the kernel returns from the current interrupt, exception, or trap. Although this is somewhat unrealistic, it is possible to configure real hardware and the kernel to behave this way, and it greatly simplifies the programming of the kernel for the project.

2.4.2. Hardware-Defined Exception Stack Frame Structure

When an interrupt, exception, or trap occurs, the hardware saves the current CPU state on the stack in an `ExceptionStackFrame` structure, and passes a pointer to this saved structure to the corresponding handler function. Each interrupt, exception, or trap handler function thus should be defined as follows:

```
void name(ExceptionStackFrame *)
```

When the handler function returns, the hardware restores the CPU state from the values then in this `ExceptionStackFrame` structure.

The format of the `ExceptionStackFrame` structure is defined in `comp421/hardware.h`. You can include the file `comp421/hardware.h` in your source file by using

```
#include <comp421/hardware.h>
```

The following fields are defined within a `ExceptionStackFrame`:

- `int vector`: The type of interrupt, exception, or trap, such as `TRAP_ILLEGAL`. This is the index into the interrupt vector table used to call this handler function.
- `int code`: A code value giving more information on the particular interrupt, exception, or trap. The meaning of the `code` value varies depending on the type of interrupt, exception, or trap. The defined values of `code` are summarized in Table 2.
- `void *addr`: This field is only meaningful for a `TRAP_MEMORY` exception. It contains the memory address whose reference caused the exception.
- `unsigned long psr`: The value of the CPU's processor status register (PSR) at the time of the interrupt, exception, or trap. The `PSR_MODE` bit in the PSR defines the current execution mode of the CPU. Specifically, if `(psr & PSR_MODE)` is true, the CPU was executing in kernel mode when the interrupt, exception, or trap occurred; if this bit is set in the `psr` value, the CPU will execute in kernel mode after returning from the handler.
- `void *pc`: The hardware program counter register value at the time of the interrupt, exception, or trap. The value of `pc` will be restored to the hardware program counter register on return from the handler.
- `void *sp`: The hardware stack pointer register value at the time of the interrupt, exception, or trap. The value of `sp` will be restored to the hardware stack pointer register on return from the handler.
- `unsigned long regs[NUM_REGS]`: The contents of eight general purpose hardware registers in the CPU at the time of the interrupt, exception, or trap (the value `NUM_REGS` is defined to be 8). For example, by convention, for a kernel call, these registers are loaded with the arguments for the call before executing the trap instruction, and thus for a `TRAP_KERNEL`, the arguments to the kernel call can be accessed using this `regs` array. More information on the use of the `regs` array in handling a kernel call is available in Section 3.2. The `regs` values will be restored to the corresponding hardware registers on return from the handler.

The values of the privileged machine registers (the `REG_*` registers, defined above in Section 2.1.2) are *not* included in the exception stack frame structure. These values are associated with the current process by the kernel, not by the hardware, and must be changed by your kernel on a context switch when/if needed.

Table 2 Meaning of Code Values in ExceptionStackFrame Structure

Vector	Code
TRAP_KERNEL	Kernel call number
TRAP_CLOCK	Undefined
TRAP_ILLEGAL	Type of illegal instruction
TRAP_MEMORY	Type of disallowed memory access
TRAP_MATH	Type of math error
TRAP_TTY_TRANSMIT	Terminal number causing interrupt
TRAP_TTY_RECEIVE	Terminal number causing interrupt

2.5. Additional CPU Machine Instructions

In addition to the hardware features of the RCS 421 computer system already described, the CPU provides two additional instructions for special purposes:

- `void Pause(void)`

The `Pause` instruction temporarily stops the CPU until the next interrupt occurs. Normally, the CPU continues to execute instructions even when there is no useful work available for it to do. In a real operating system on real hardware, this is alright since there is nothing else for the CPU to do. Operating Systems usually provide an idle process which is executed in this situation, and which is typically an empty infinite loop. However, in order to be nice to other users on the real Unix machines you will be using in the project, your idle process should not just be an empty loop in this way. Instead, the idle process in your kernel should be a loop that executes the `Pause` instruction in each loop iteration.

- `void Halt(void)`

The `Halt` instruction completely stops the CPU (end ends the execution of the simulated RCS 421 computer system). By executing the `Halt` hardware instruction, the CPU is completely halted and does not begin execution again until rebooted (i.e., until started again by running your kernel again from the shell on your Unix terminal).

3. Yalnix Operating System Details

This section describes the Yalnix operating system, a reasonably simple, mythical operating system that can be implemented on many types of computer systems including the RCS 421.

This section is written partially in the style of a real operating system architecture manual. Everything in this section is something that a real designer, implementor, or user of the operating system would need to know about the operating system. Certain features of the RCS 421 computer system on which your kernel will execute are mentioned here only because you will need them in the project; the Yalnix operating system could also be implemented on other hardware platforms.

3.1. Yalnix Kernel Calls

User processes request services from the operating system kernel by executing a “kernel call” (sometimes called a “system call” or “syscall”). The following kernel calls are defined in Yalnix:

- `int Fork(void)`

`Fork` is the only way a process can create a new process in Yalnix. The memory image of the new process (the child) is a copy of that of the process calling `Fork` (the parent). When the `Fork` call completes, *both* the parent process and the child process return (separately) from the `Fork` kernel call as if each had been the one to call `Fork`, since the child is a copy of the parent. The only distinction is the fact that the return value in the calling (parent) process is the (nonzero) process ID of the new (child) process, while the value returned in the child is 0. If, for any reason, the new process cannot be created, `Fork` instead returns the value `ERROR` to the calling process.

- `int Exec(char *filename, char **argvec)`

Replace the currently running program in the calling process's memory with the program stored in the file named by `filename`. The argument `argvec` points to a vector of arguments to pass to the new program as its argument list. The new program is called as

```
main(int argc, char **argv)
```

where `argc` is the argument count and `argv` is an array of character pointers to the arguments themselves. The strings pointed to by the entries in `argv` are copied from the strings pointed to by the `argvec` array passed to `Exec`, and `argc` is a count of entries in this array before the first `NULL` entry, which terminates the argument list. When the new program begins running, its `argv[argc]` is `NULL`. By convention the first argument in the argument list passed to a new program (`argvec[0]`) is also the name of the new program to be run, but this is just a convention; the actual file name to run is determined only by the `filename` argument. On success, there is no return from this call in the calling program, and instead, the new program begins executing in this process at its entry point, and its `main(argc, argv)` routine is called as indicated above. On failure, this call returns `ERROR` to the calling program as the result of `Exec` and does not run the new program.

- `void Exit(int status)`

`Exit` is the normal means of terminating a process. The current process is terminated, and the integer `status` value is saved for possible later collection by the parent process on a call to `Wait`. All resources used by the calling process are freed, except for the saved `status` information. This call can never return. When a process exits or is terminated by the kernel, if it has children, they should continue to run normally, but they will no longer have a parent. When such an orphan process later exits, its exit status is not saved or reported to its parent process since it no longer has a parent process.

- `int Wait(int *status_ptr)`

Collect the process ID and exit status returned by a child process of the calling program. When a child process `Exits`, its exit status information is added to a FIFO queue of child processes not yet collected by its specific parent. After the `Wait` call, this child process information is removed from the queue. If the calling process has no remaining child processes (exited or running), `ERROR` is returned and the integer pointed to by `status_ptr` is not modified. Otherwise, if there are no exited child processes waiting for collection by this calling process, the calling process is blocked until its next child calls `exits` or is terminated by the kernel. On success, the process ID of the child process is returned and its exit status is copied to the integer referenced by the `status_ptr` argument. On any error, this call instead returns `ERROR`.

- `int GetPid(void)`

Returns the process ID of the calling process.

- `int Brk(void *addr)`

`Brk` sets the operating system's idea of the lowest location not used by the program (called the "break") to `addr`. Since memory is allocated to a process in units of whole pages, the actual break is in effect rounded up to a multiple of `PAGESIZE`, although the specified `addr` value need not itself be a multiple of `PAGESIZE`. This call has the effect of allocating or deallocating enough memory to cover only up to the specified address. Locations not less than the break (and below the allocated memory for the process's stack) are not in the address space of the process and will thus cause an exception if accessed. The value 0 is returned by `Brk` on success. If any error is encountered (for example, if not enough memory is available or if the address `addr` is invalid), the value `ERROR` is returned.

- `int Delay(int clock_ticks)`

The calling process is blocked until `clock_ticks` clock interrupts have occurred after the call. Upon completion of the delay, the value 0 is returned. If `clock_ticks` is 0, return is immediate. If `clock_ticks` is less than 0, `ERROR` is returned instead.

- `int TtyRead(int tty_id, void *buf, int len)`

Read the next line of input (or a portion of it) from terminal `tty_id`, copying the bytes of input into the buffer referenced by `buf`. The maximum length of the line to be returned is given by `len`. The line returned in the buffer is *not* null-terminated.

The calling process is blocked until a line of input is available to be returned. If the length of the next available input line is longer than `len` bytes, only the first `len` bytes of the line are copied to the calling process's buffer, and the remaining bytes of the line are saved by the kernel for the next `TtyRead` (by this or another process) for this terminal. If the length of the next available input line is shorter than `len` bytes, only as many bytes are copied to the calling process's buffer as are available in the input line. On success, the number of bytes actually copied into the calling process's buffer is returned; in case of any error, the value `ERROR` is returned.

- `int TtyWrite(int tty_id, void *buf, int len)`

Write the contents of the buffer referenced by `buf` to the terminal `tty_id`. The length of the buffer in bytes is given by `len`. The calling process is blocked until all characters from the buffer have been written on the terminal. On success, the number of bytes written (which should be `len`) is returned; in case of any error, the value `ERROR` is returned.

Function prototypes for all of the Yalnx kernel calls defined above are defined in the include file `comp421/yalnx.h`, in the directory `/clear/courses/comp421/pub/include/comp421`. You can include the file `comp421/yalnx.h` in your source file by using

```
#include <comp421/yalnx.h>
```

Some additional Yalnx kernel calls are also defined in `comp421/yalnx.h`, although you need not implement them for this project.

Unless otherwise noted above, all Yalnx kernel calls return the value `ERROR` in case of any error in processing that kernel call. The value `ERROR` is defined in the include file `comp421/yalnx.h`.

The file `comp421/yalnx.h` also defines a constant "kernel call number" each type of kernel call, used to indicate to the kernel the specific kernel call being invoked; this type value is visible to the kernel in the `code` field of the `ExceptionStackFrame` structure passed by the hardware for the trap, as described in Section 3.2.

Your kernel must verify all arguments passed to a kernel call, and should return `ERROR` if any arguments are invalid. In particular, this includes checking that a pointer value passed to the kernel call is valid before you use the pointer in your kernel. Specifically, your kernel must verify that the pointer points into the process's memory and should look in the page table to make sure that the *entire* area (such as defined by a pointer and a specified length) is readable and/or writable (as appropriate) before the kernel actually tries to read or write there. For C-style character strings (null-terminated) you will need to, in effect, check the pointer to each byte as you go (C strings like this are passed to `Exec`). For simplicity, you should write a common routine to check a buffer with a specified pointer and length for read, write and/or read/write access; and a separate routine to verify a string pointer for read access. The string verify routine would, in effect, check access to each byte, checking each until it found the `'\0'` at the end. Insert calls to these two routines as needed at the top of each kernel call to verify the pointer arguments before you use them.

Such checking of arguments is important for security and reliability. An unchecked `TtyRead`, for instance, might well overwrite crucial parts of the operating system, which might in some clever way gain an intruder access as a privileged user. Also, a pointer to memory that is not correctly mapped or not mapped at all would generate a `TRAP_MEMORY`, causing the kernel to crash.

3.2. Interrupt, Exception, and Trap Handling

As described above in Section 2.4, each type of interrupt, exception, or trap that can be generated by the hardware has a corresponding type value used by the hardware to index into the interrupt vector table in order to find the address of the handler function to call. The operating system kernel must create the interrupt vector table and initialize each corresponding entry in the table with a pointer to the relevant handler function within the kernel. The address of the interrupt vector table must also be written into the `REG_VECTOR_BASE` register by the kernel at boot time.

When an interrupt, exception, or trap occurs, the hardware automatically saves the current CPU state on the kernel stack in a `ExceptionStackFrame` data structure. When calling the handler function, the hardware passes a pointer to this `ExceptionStackFrame` structure as the procedure argument to the handler.

The range of possible index values in the interrupt vector table range from 0 to `TRAP_VECTOR_SIZE`, and the interrupt vector table must be exactly this size. Each entry not needed must be initialized to `NULL`.

The defined kernel call type values, together with an overview of the operation that should be performed by the kernel in response to each, are listed below:

- `TRAP_KERNEL`: Execute the requested kernel call, as indicated by the kernel call number in the `code` field of the `ExceptionStackFrame` passed by reference to this trap handler function. The arguments passed by the user process for this kernel call are available to the kernel beginning in `regs[1]` in the `ExceptionStackFrame` passed by reference to the kernel's `TRAP_KERNEL` handler function. The return value from the kernel call should be returned to the user process in the `regs[0]` field of the `ExceptionStackFrame`.
- `TRAP_CLOCK`: The Yalnx kernel should implement round-robin process scheduling with a CPU quantum per process of 2 clock ticks. After the current process has been running as the current process continuously for at least 2 clock ticks, if there are other runnable processes on the ready queue, perform a context switch to the next runnable process.
- `TRAP_ILLEGAL`: Terminate the currently running Yalnx user process and print a message giving the process id of the process and an explanation of the problem; and continue running other processes.
- `TRAP_MEMORY`: The kernel must determine if this exception represents an implicit request by the current process to enlarge the amount of memory allocated to the process's stack, as described in

more detail in Section 3.4.4. If so, the kernel enlarges the process's stack to "cover" the address that was being referenced that caused the exception (the `addr` field in the `ExceptionStackFrame`) and then returns from the exception, allowing the process to continue execution with the larger stack. Otherwise, terminate the currently running Yalrix user process and print a message giving the process id of the process and an explanation of the problem; and continue running other processes.

- `TRAP_MATH`: Terminate the currently running Yalrix user process and print a message giving the process id of the process and an explanation of the problem; and continue running other processes.
- `TRAP_TTY_RECEIVE`: This interrupt signifies that a new line of input is available from the terminal indicated by the `code` field in the `ExceptionStackFrame` passed by reference to this interrupt handler function. The kernel should read the input from the terminal using a `TtyReceive` hardware operation and if necessary should buffer the input line for a subsequent `TtyRead` kernel call by some user process.
- `TRAP_TTY_TRANSMIT`: This interrupt signifies that the previous `TtyTransmit` hardware operation on some terminal has completed. The specific terminal is indicated by the `code` field in the `ExceptionStackFrame` passed by reference to this interrupt handler function. The kernel should complete the `TtyWrite` kernel call and unblock the blocked process that started this `TtyWrite` kernel call that started this output. If other `TtyWrite` calls are pending for this terminal, also start the next one using `TtyTransmit`.

As described above, a `TRAP_KERNEL` trap occurs when a Yalrix user process requests a kernel call for some function provided by the kernel. A user processes calls the kernel by executing a trap instruction. However, since the C compiler cannot directly generate a trap instruction in the generated machine code for a program, a library of assembly language routines is provided that performs this trap for the user process. This library provides a standard C procedure call interface for kernel calls, as indicated in the description of each kernel call in Section 3.1. The trap instruction generates a trap to the hardware, which invokes the kernel using the `TRAP_KERNEL` vector from the interrupt vector table.

Upon entry to your kernel's `TRAP_KERNEL` handler function, the `code` field of the `ExceptionStackFrame` structure indicates which kernel call is being invoked (as defined with symbolic constants in `comp421/yalrix.h`). You can include the file `comp421/yalrix.h` in your source file by using

```
#include <comp421/yalrix.h>
```

The arguments to this call, supplied to the library procedure call interface in C/C++, are available in the `regs` array in the `ExceptionStackFrame` structure received by your `TRAP_KERNEL` handler. Each argument passed to the library procedure is available in a separate `regs` entry, in the order passed to the procedure, beginning with `regs[1]`. In other words, the first argument to the kernel call is in `regs[1]`, the second argument to the kernel call is in `regs[2]`, and so on.

Each kernel call returns a single integer value, indicated in Section 3.1, which becomes the return value from the C library procedure call interface for the kernel call. When returning from a `TRAP_KERNEL` trap, the value to be returned from this interface should be placed by the kernel in `regs[0]` in the `ExceptionStackFrame`.

When the kernel terminates a process (e.g., after a `TRAP_ILLEGAL`, not as a result of `Exit()`), your kernel should print a message, giving the process id of the process and some explanation of the problem. The exit status reported to the parent process of the terminated process when the parent calls the `Wait` kernel call (as described in Section 3.1) should be the value `ERROR`. You can print the message from your kernel with `printf` or `fprintf`, in which case it will print out on the Unix window from which you

started your kernel, or (better, but not required) you can use the hardware-provided `TtyTransmit` to print the message on the Yalrix console.

3.3. Kernel Booting Entry Point

Your kernel is not a complete program—it is more like a library of functions that get called on interrupts, exceptions, and traps. As such, your kernel does not have a `main` procedure like typical C or C++ programs.

Instead, when booted, your kernel begins executing at a procedure named `KernelStart`, which you must write in your kernel. The procedure named `KernelStart` is automatically called by the bootstrap firmware in the computer, as follows:

```
void KernelStart(ExceptionStackFrame *frame,
                 unsigned int pmem_size, void *orig_brk, char **cmd_args)
```

The procedure arguments passed to `KernelStart` are built by the bootstrap firmware during initialization of the machine. The `KernelStart` procedure should initialize your operating system kernel and then return. The arguments to `KernelStart` are defined as follows:

- The `frame` argument is a pointer to an initial `ExceptionStackFrame` structure (defined in Section 2.4.2). The `KernelStart` procedure is called in the same way as the interrupt, exception, and trap handlers, as described in Sections 2.4 and 3.2, except that `KernelStart` is passed additional arguments (described below). This initial `ExceptionStackFrame` records the state of the machine at boot time, and any changes made in the values in this `ExceptionStackFrame` control how and where the machine will execute when `KernelStart` returns. The `KernelStart` procedure should initialize the operating system kernel and then return, starting the first process executing. More details on initializing the operating system kernel and in the use of the `frame` argument passed to `KernelStart` are provided in Section 3.6.
- The `pmem_size` argument gives the size of the physical memory of the machine you are running on, as determined dynamically from the hardware by the bootstrap firmware. The size of physical memory is given in units of bytes.
- The `orig_brk` argument gives the initial value of the kernel’s “break.” That is, this address is the first address that is *not* part of the kernel’s initial heap.
- The `cmd_args` argument is a vector of strings (in the same format as `argv` for normal Unix `main` programs), containing a pointer to each argument from the boot command line (what you typed at your Unix terminal) to start the machine and thus the kernel. The `cmd_args` vector is terminated by a `NULL` pointer.

You must write a procedure named `KernelStart` as part of your kernel. Since this is the place at which your kernel first starts executing at boot time, it is recommended that you write and test a simple version of a `KernelStart` procedure as the first step in building your kernel.

3.4. Memory Management

Your kernel is responsible for all aspects of memory management, both for user processes executing on the system and for the kernel’s own use of memory.

Memory management for user processes is reasonably straightforward. To enlarge the program area (the heap) of the user process, the process calls the `Brk` kernel call. In general, processes don’t do this

explicitly themselves, however, since, for example, the standard C library `malloc` function will call `Brk` automatically if it needs more memory. To instead enlarge the stack area of the user process, the process simply tries to use the new stack pointer value, which causes a `TRAP_MEMORY` exception if not enough memory is currently allocated for the process's stack. Details of the procedure for growing the stack of a user process are given in Section 3.4.4.

Memory management for the kernel itself is a bit more complicated. The kernel may also use `malloc` for dynamic memory allocation for its own data structures, but unlike `malloc` from a user program, the kernel can't rely on someone else to actually allocate the memory for it. The kernel must allocate the physical memory into Region 1 for the kernel's own use, as needed for `malloc` calls made by the kernel. Details of this procedure are given in Section 3.4.2. In addition, the kernel must initialize and enable virtual memory at boot time; details of this procedure are given in Section 3.4.3.

3.4.1. Virtual Address Space Layout

The Yalnix operating system uses the virtual memory management hardware provided by the RCS 421 computer system to implement a separate virtual address space for each process, with the kernel shared by all processes. The layout of the virtual address space implemented by Yalnix is shown in Figure 3.

The machine's virtual address space is divided into two "Regions," named Region 0 and Region 1. Region 0 is used by each user process, and Region 1 is used by the kernel and is shared by all processes. Each of the two regions is the same size, defined as `VMEM_REGION_SIZE` (which also equals `VMEM_0_SIZE` and `VMEM_1_SIZE`). Region 0 begins at virtual address `VMEM_0_BASE` (which is defined equal to 0) and ends at virtual address `VMEM_0_LIMIT`. Region 1 begins at virtual address `VMEM_1_BASE` (which is defined equal to `VMEM_0_LIMIT`) and ends at virtual address `VMEM_1_LIMIT`. The constant `VMEM_BASE` (which is defined equal to `VMEM_0_BASE` and is thus equal to 0) is the virtual address of the beginning of the entire virtual address space, and the constant `VMEM_LIMIT` (which is defined equal to `VMEM_1_LIMIT`) is the virtual address of the end of the entire virtual address space.

The limit of the program's heap (the lowest address above the heap that is not part of the process' address space) is called the process's *break* and is set from Yalnix user processes by the `Brk` kernel call, which you will implement in your kernel.

The break of a user process is initialized (when a program begins running) to an address defined in the executable program file from which the process was loaded by `Exec`. Details on how to load a program for `Exec` are provided in the file `load.template`, which is a template for a function `LoadProgram` that can load a Yalnix program into memory. The file `load.template` is located in the directory `/clear/courses/comp421/pub/templates`. You should copy and modify this template as described in the `load.template` file itself. After the `Exec` completes and a program begins execution, the process itself allocates more memory as needed by calling the kernel's `Brk` kernel call.

The user stack for each process begins (at the top of the stack) at virtual address `USER_STACK_LIMIT`. The size of the user stack of a process will be automatically grown by your Yalnix kernel in response to a `TRAP_MEMORY` exception resulting from an access to a page between the current break and the current stack pointer that has not yet been allocated. This happens during normal execution as the program makes procedure calls that automatically expand the stack downwards. Details of this are given in Section 3.4.4 below.

Each process executing on Yalnix also has a separate kernel stack, located in that process's Region 0. The kernel stack is fixed in size at `KERNEL_STACK_PAGES` pages (or `KERNEL_STACK_SIZE` bytes). The kernel stack begins at (at the top of the stack) at virtual address `KERNEL_STACK_LIMIT` and grows down (to lower virtual addresses) from there. The bottom address limit of the kernel stack is virtual address `KERNEL_STACK_BASE`. Each time the hardware calls a handler procedure in the kernel, the hardware pushes the `ExceptionStackFrame` information onto the kernel stack, which includes the

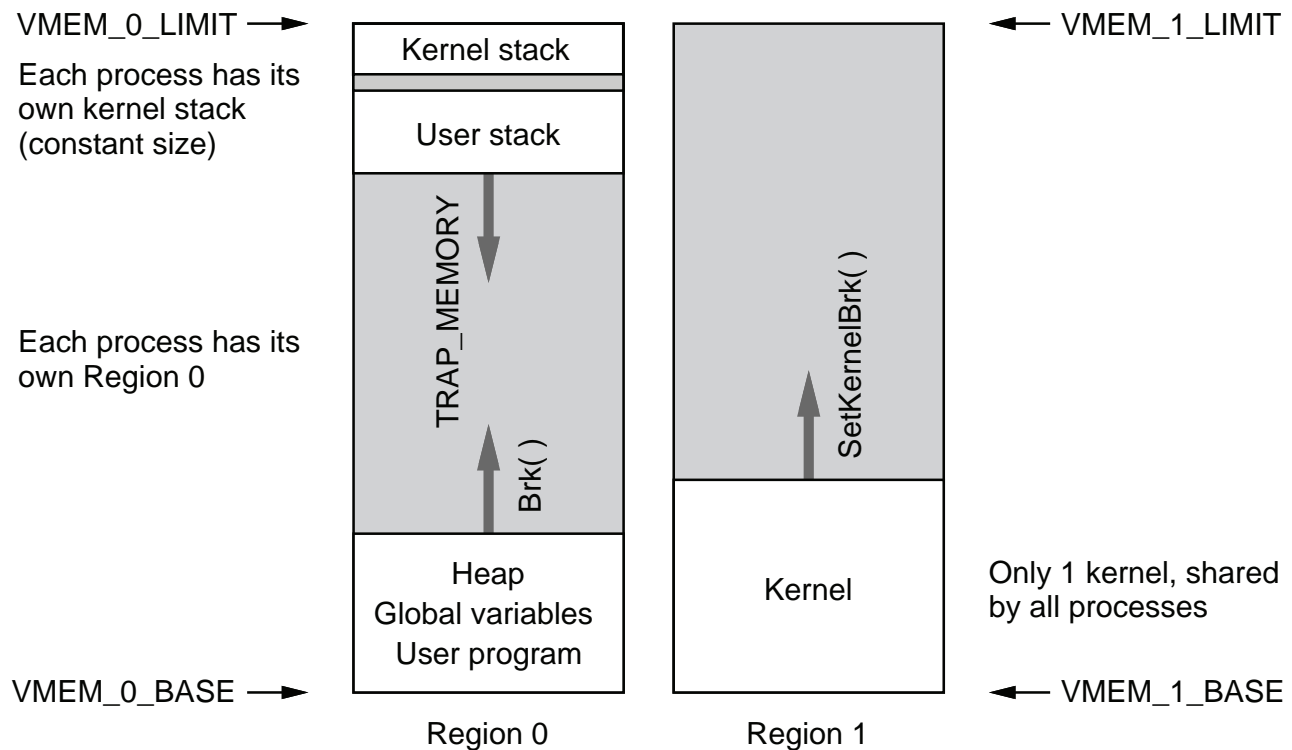


Figure 3 Yalrix Virtual Address Space Layout

then-current value of the stack pointer register. Before pushing this information onto the stack, the hardware automatically changes the stack pointer to the value `KERNEL_STACK_LIMIT`, and the value of the stack pointer register saved in the `ExceptionStackFrame` is the original value of the stack pointer before the value of this register is changed to `KERNEL_STACK_LIMIT`.

The kernel stack is located in Region 0, and there is a separate kernel stack for each process. By having a separate kernel stack for each process, it is possible for the kernel to block the current process while inside the kernel. When performing a context switch, the kernel can then switch to the kernel stack of the newly selected process (being context switched in), thus leaving the state of the old process's kernel stack entirely unmodified until that process is selected to run again next and is context switched in. All of that process's local variables inside the kernel and all of the subroutine return addresses and such that represent the dynamic state of the process inside the kernel when it was last context switched out, are all still intact and again available when the process is next context switched in. The process can thus begin execution again after the context switch at exactly the point it was at inside the kernel when it last ran.

3.4.2. Kernel Memory Management

Unlike the user stack for each user process, the kernel stack for each process is fixed in size, as noted above in Section 3.4.1. Thus, after setting up the kernel stack for a process, the kernel need not explicitly manage the kernel stack for that process.

In contrast to the kernel stack, the kernel heap *does* require management by your operating system kernel. Each time some procedure such as `malloc` called by your kernel needs more pages added to the kernel's memory, it calls a function named `SetKernelBrk`, defined as follows:

```
int SetKernelBrk(void *addr)
```

In a *user* program, `malloc` calls the `Brk` kernel call to request the operating system kernel to add additional memory to the process's address space. When the kernel calls `malloc`, however, the internal kernel function `SetKernelBrk` is called instead. (Internally, `malloc` keeps a list of free block of memory and carves up existing free blocks as necessary, and only asks for more total memory kernel when its existing available memory cannot satisfy the requested allocation.)

Before your kernel has enabled virtual memory, the entire physical memory of the machine is already available to your kernel. The job of `SetKernelBrk` in this case is very easy. You need only keep track of the location of the current break for the kernel. The initial location of the break is passed to your kernel's `KernelStart` procedure as the `orig_brk` argument, as described in Section 3.3. Calls to `SetKernelBrk`, before virtual memory has been enabled, simply move this break location up (or down) to the new location specified as `addr` in the `SetKernelBrk` call. You will need to know the current location of your kernel's break in order to correctly initialize page table entries for your kernel's Region 1 before enabling virtual memory.

After you have enabled virtual memory, the job of the `SetKernelBrk` function that you will write becomes more complicated. After this point, your implementation of `SetKernelBrk`, when called with an argument of `addr`, must allocate physical memory page frames and map or unmap virtual pages as necessary to make `addr` the new kernel break. In your kernel, you should keep a flag to indicate if you have yet enabled virtual memory, so your implementation of `SetKernelBrk` can know which behavior (for before virtual memory has been enabled, or for after virtual memory has been enabled) it should perform for each call to `SetKernelBrk`.

Should your kernel run out of physical memory or otherwise be unable to satisfy the requirements of `SetKernelBrk`, your `SetKernelBrk` function can return `-1`. This will eventually cause `malloc` in the kernel to correctly return a `NULL` pointer, as is defined for the standard behavior of `malloc`. If `SetKernelBrk` is successful (the likely case) it should return `0`.

Note that the `malloc` implementation that you use inside your kernel (and its helping function `SetKernelBrk`, which you must write) is completely separate from the `malloc` implementation that each user process uses. Thus, the kernel has its own break location and the list of free blocks of memory managed by `malloc` is completely separate from that in any process. Likewise, each process running on Yalnx has its own break and is linked with its own copy of the `malloc` library routines (if the process uses dynamic memory allocation), making its list of free blocks of memory also separate from each other process.

3.4.3. Initializing the Page Tables and Virtual Memory

As described in Section 2.2.6, when the machine is booted, the machine's virtual memory subsystem is initially disabled and is only enabled once your operating system kernel writes a nonzero value into the `REG_VM_ENABLE` register. Before doing this, your kernel must initialize the page table and page table registers.

A consequence of this procedure for enabling virtual memory is that the kernel will have already been loaded into memory and executed for a while before virtual memory is enabled. During this execution, all addresses used by the kernel during its execution are treated by the hardware as *physical* addresses, but after enabling virtual memory, all addresses then used by the kernel during its execution are treated by the hardware as *virtual* addresses. This means that, at the point at which virtual memory is first enabled, suddenly all addresses are interpreted differently. Without special care in the kernel, this can cause great confusion to the kernel, since by this time, many addresses are in use in the running kernel: for example, in the program counter and stack pointer, perhaps in some of the general registers, and in any pointer variables declared and used in the kernel.

To solve this problem, the simplest method is to arrange the initial page table entries for the kernel so that all addresses actually still point to the same places after enabling virtual memory. This can be done by the kernel by arranging for the initial virtual address space layout to be the same as the current physical address space layout, as illustrated in Figure 4. The hardware makes this possible by guaranteeing that `PMEM_BASE`, the beginning address of physical memory, is the same value as `VMEM_BASE`, the beginning address of virtual memory (which is thus also the beginning address of Region 0 of virtual memory). In building the initial page tables in the kernel before enabling virtual memory, the kernel should make sure, for every page of physical memory `pfn` then in use by your kernel, that a page table entry is built so that the new virtual page number `vpn = pfn` maps to physical page number `pfn`.

To correctly build these initial page table entries, your kernel must know how many page table entries to build, and it must know the correct protection settings for each of these page table entries. In particular, the text page table entries for your kernel should be set to “read” and “execute” protection for kernel mode, and the data/bss/heap page table entries for your kernel should be set to “read” and “write” protection for kernel mode; the user mode protection for both kinds of kernel page table entries should be “none” (no access).

The end of the kernel’s heap can be known by your kernel from the initial break address passed as `orig_brk` to your `KernelStart` function and from any calls to your `SetKernelBrk` function that may have occurred so far in the execution of your kernel. The end of the kernel’s text can be known by your kernel from the following value:

`&_etext`

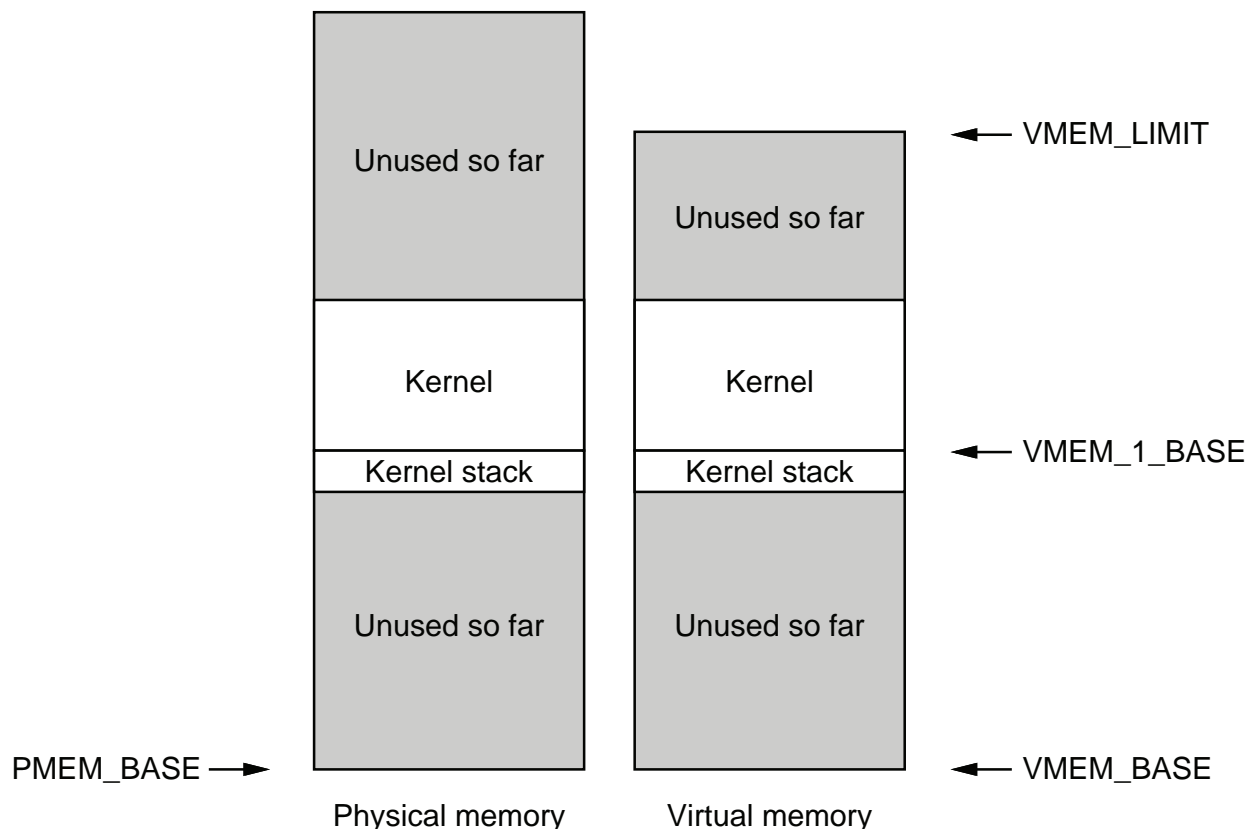


Figure 4 Initial Memory Layout Before and After Enabling Virtual Memory

The value of `&_etext` is the lowest address that does not contain executable instructions (text) in your kernel. The symbol `_etext` is placed in the symbol table of your kernel by the linker, although it is not really a variable. The *address* of this symbol is what you want.

3.4.4. Growing a User Process's Stack

When a process pushes onto its user mode stack, it simply decrements the current stack pointer value and attempts to write into memory at the location now pointed to by the stack pointer. Normally, this works with no problem, since this same memory was likely used to store other data earlier pushed onto and popped off of the stack. However, at times, the stack will grow larger than it has been before in this process, growing into a new page, which will cause a `TRAP_MEMORY` to be generated by the hardware when the process tries to write to the memory pointed to by its new stack pointer value. As noted in Section 3.2, the `TRAP_MEMORY` in this case should be interpreted by your kernel as an implicit request to grow the process's user stack.

In your handler for a `TRAP_MEMORY` exception, if the virtual address being referenced that caused the exception (the `addr` value in the `ExceptionStackFrame`) is in Region 0, is below the currently allocated memory for the stack, and is above the current break for the executing process, your Yalnx kernel should attempt to grow the stack to “cover” this address, if possible. In all other cases, you should terminate the currently running Yalnx user process, but should continue running any other user processes.

When growing the stack, you should leave at least one page unmapped (with the valid bit in its page table set to 0) between the heap and user stack in Region 0, so the stack will not silently grow into and overlap with the heap without triggering a `TRAP_MEMORY`. This unmapped page between the program area and the stack is known as a “red zone” into which the stack (and the break) should not be allowed to grow. You must also check that you have enough physical memory to grow the stack.

Note that a user-level procedure call with many large local variables may grow the stack by more than one page in one step. The unmapped page that is used to form a red zone for the stack as described in the previous paragraph is therefore not a guaranteed reliable safety measure, but it does add some assurance that the stack will not grown into the program.

3.5. Context Switching

When some handler function in the kernel is called by the hardware due to an interrupt, exception, or trap, the hardware enters kernel mode and switches to the kernel stack. The current user process stack then contains the local variables and return addresses from each function called by the user process leading up to the point at which the interrupt, exception, or trap occurred. As the kernel's handler function executes (and possibly calls other internal functions inside the kernel), the kernel stack then contains the local variables and return addresses for each function inside the kernel called by this process during the execution since entering the kernel, including the handler function itself. Both of these stacks (the user stack and the kernel stack) are located in virtual memory in Region 0. Also in Region 0 is the text, data, bss, and heap used by the executing user process.

A context switch from one process to another can only be performed while executing inside the kernel. To perform a context switch, the kernel switches from one Region 0 page table to another Region 0 page table (and also, at the same time, writes to the `REG_TLB_FLUSH` register to flush all the Region 0 entries from the TLB). This causes the kernel stack, as well as the user stack and the rest of the contents of Region 0 to switch. However, this switch occurs at some place deep inside your kernel, at whatever point in the code your kernel performed the necessary operations to accomplish this switch. Since your kernel would then be executing on the kernel stack, switching the kernel stack while using the kernel stack could create great confusion for the kernel's own execution. For example, the return address and local variables

for the current function would suddenly “disappear” when the Region 0 virtual memory is replaced with a different Region 0.

To simplify the operation of context switching and to avoid this confusion to the kernel’s own execution, we provide the function

```
int ContextSwitch(SwitchFunc_t *func, SavedContext *ctxp,
                 void *p1, void *p2)
```

to help with the context switch process. The type `SwitchFunc_t` (switch function type) is a messy C typedef of a special kind of function. `ContextSwitch` and the type `SwitchFunc_t` are defined in `comp421/hardware.h`, in the directory `/clear/courses/comp421/pub/include/comp421`. The `ContextSwitch` function temporarily stops using the standard kernel stack and calls the function `func` using a temporary, different stack not located in Region 0.

For example, suppose you named your switch function `MySwitchFunc`. Then you would declare `MySwitchFunc` as:

```
SavedContext *MySwitchFunc(SavedContext *ctxp, void *p1, void *p2)
```

Since `MySwitchFunc` is called while not using the normal kernel stack, `MySwitchFunc` can easily and safely do what it needs to do to switch Region 0 to complete the context switch.

The two “void *” arguments `p1` and `p2` passed to `ContextSwitch` will be passed unmodified to `MySwitchFunc`. You should use them, for example, to point to the current process’s PCB and to the PCB of the new process to be context switched in, respectively. In `MySwitchFunc`, you should do whatever you need to do to context switch between these two processes.

The `ctxp` argument (the pointer itself) passed to `ContextSwitch` is likewise passed unmodified to `MySwitchFunc`. However, before calling `MySwitchFunc`, the `ContextSwitch` function saves a copy of the current CPU state (e.g., registers) in this `SavedContext` structure. When returning from `MySwitchFunc`, you should return a pointer to the `SavedContext` that records the CPU state you want to have restored when `ContextSwitch` returns. Inside your function `MySwitchFunc`, before returning, you should also do the necessary operations to switch from one process’s Region 0 page table to the other process’s Region 0 page table.

Specifically, you should include a `SavedContext` within each process’s PCB. Suppose this member of the PCB structure is called `ctx`, and suppose `pcb1` currently points to the PCB of the process that you are context switching from and `pcb2` points to the PCB of the process you are switching to. Then, to context from a first process indicated by `pcb1` to a second process indicated by `pcb2`, you would call `ContextSwitch` as follows:

```
ContextSwitch(MySwitchFunc, &pcb1->ctx, (void *)pcb1, (void *)pcb2)
```

Inside `MySwitchFunc`, you should switch from using the Region 0 page table associated with the process indicated by `pcb1` to using the Region 0 page table associated with the process indicated by `pcb2`. Then, when `MySwitchFunc` returns, it should return `&pcb2->ctx`. To later context switch back to the first process, you would call `ContextSwitch` again with the roles of `pcb1` and `pcb2` reversed. The switching from one Region 0 page table to another *must* be done from within `MySwitchFunc`.

The actual `SavedContext` structure definition is provided by `comp421/hardware.h`. However, you don’t need to worry what is in a `SavedContext`. On a real machine, this structure would be full of low-level hardware-specific stuff; on the RCS 421 simulation, it is full of low-level Unix-specific (and internal simulation-specific) stuff.

When you call `ContextSwitch`, the current process will be blocked inside the kernel exactly where it called `ContextSwitch`, with exactly the state that it had at that time. `ContextSwitch` and

`MySwitchFunc` do the context switch, and when this process is later context switched back to, this process's earlier call to `ContextSwitch` returns. The return value of `ContextSwitch` in this case is 0. If, instead, any error occurs, `ContextSwitch` does not switch contexts and instead returns -1; in this case, you should print an error message and exit, as this generally indicates a programming error in your kernel.

Note that `ContextSwitch` doesn't do anything to move a PCB from one queue to another in your kernel, or to do any bookkeeping as to why that process was context switched out or when it can be context switched in again. `ContextSwitch` only helps you with actually switching the state of the process and later restoring it. Your kernel has to take care of the rest itself.

Depending on what you make your `MySwitchFunc` (or any other name you choose) function do, you can also use `ContextSwitch` to just get a copy of the current `SavedContext` and kernel stack, without necessarily switching to a new process. You might find this useful in your implementation of `KernelStart` in creating the *idle* or *init* process, and in your implementation of the `Fork` kernel call in creating the new child process.

3.6. Bootstrapping and Kernel Initialization

As noted above in Section 3.3, your kernel begins execution at a function named `KernelStart`, which you must write. Your kernel does *not* have a main function. Your `KernelStart` function is called with a pointer to an `ExceptionStackFrame` structure (among other arguments). Your kernel should use this `ExceptionStackFrame` as the basis for “cloning” others. In particular, your `KernelStart` function should create both an *idle* and an *init* process. One of these (presumably *init*) will run as the first process when `KernelStart` returns. The other of these two processes will be context switched to possibly sometime later. At that time, this process must have an initialized kernel stack, since when `ContextSwitch` returns after this context switch, this process will be executing in your kernel. Likewise, you must have a corresponding initialized `SavedContext` to use in loading the CPU state of this process.

In general, before allowing the execution of user processes, the `KernelStart` routine should perform any initialization necessary for your kernel or required by the hardware. In particular, in addition to any initialization you may need to do for your own data structures, your kernel initialization should probably include the following steps (the order of some of these steps can be changed to suite how you want to structure your initialization):

- Initialize the interrupt vector table entries for each type of interrupt, exception, or trap, by making them point to the correct handler functions in your kernel.
- Initialize the `REG_VECTOR_BASE` privileged machine register to point to your interrupt vector table.
- Build a structure to keep track of what page frames in physical memory are free. For this purpose, you might be able to use a linked list of physical frames, implemented in the frames themselves. Or you can have a separate structure, which is probably easier, though slightly less efficient. This list of free page frames should be based on the `pmem_size` argument passed to your `KernelStart`, but should of course not include any memory that is already in use by your kernel.
- Build the initial page tables for Region 0 and Region 1, and initialize the registers `REG_PTR0` and `REG_PTR1` to define these initial page tables.
- Enable virtual memory.
- Create an “idle” process to be run by your kernel when there are no other runnable (ready) processes in the system. The idle process should be a loop that executes the `Pause` machine instruction on each loop iteration.

- Create the first process and load the initial program into it. In this step, guide yourself by the file `load.template` that we provide, which shows a skeleton of the procedure necessary to load an executable program from a Unix file into memory as a Yalrix process. This initial process will serve the role of the “init” process in Unix as the parent (or grandparent, etc.) of all processes running on the system. To run your initial program you should put the file name of the init program on your shell command line when you run your kernel. This program name will then be passed to your `KernelStart` as one of the `cmd_args` strings.
- Return from your `KernelStart` routine. The machine will begin running the program defined by the current page tables and by the values returned in the `ExceptionStackFrame` structure (values which you have presumably modified to point to the program counter and stack pointer, etc., of the initial process).

3.7. Loading Yalrix User Programs from Disk

When creating the *init* process in `KernelStart` and when loading a new program in your implementation of `Exec`, you will need to open and read the new program from a Unix file from inside your kernel. Although a bit unrealistic (a real kernel would use that system’s own file system), reading the program from a Unix file allows you to use the normal Unix file I/O calls in this project, such as `open`, `read`, and `close` (or `fopen`, `fread`, and `fclose`, if you prefer).

Although this part of loading a new program can be simplified for this project, you still need to be able to understand what you read from the program file and must know how to set up the memory of a new program from the file. This can’t be simplified very much and still do the necessary job. So, we provide you a template for the code necessary to do this.

Specifically, we provide a template for a function called `LoadProgram`, which can load a program from a Unix executable file into a Yalrix process address space. This function template is in the file `load.template` in the directory `/clear/courses/comp421/pub/templates`. You should make a copy of this file in your own directory and edit it to fit with your kernel. There are a number of places in that file that you must modify, and there are instructions in the file on how to do these modifications. Look for places marked with the symbol `>>>>` and follow the instructions there. *This file will not compile correctly until you make all the necessary modifications and remove or comment out the >>>> instructions there.*

3.8. Additional Terminal Interface

As described above in Section 2.3.1, the hardware interface to the terminals is through the `TtyTransmit` and `TtyReceive` hardware instructions. Only your kernel can execute these instructions. Your Yalrix kernel must provide kernel calls `TtyWrite` and `TtyRead` to allow user programs access to the terminals.

As a convenience, we also provide a library routine that can be used by Yalrix *user* processes:

```
int TtyPrintf(int, char *, ...)
```

that works similarly to the standard Unix `printf` function, but does its output using the Yalrix `TtyWrite` kernel call. The first argument to `TtyPrintf` is the Yalrix terminal number to write to. The next argument is a standard `printf`-style C format string, and any remaining arguments are the values for that format string to format. The return value from `TtyPrintf` is the value returned by your kernel from the underlying `TtyWrite` call performed by `TtyPrintf`. The maximum length of formatted output that can be written by a single call to `TtyPrintf` is `TERMINAL_MAX_LINE` bytes (defined in `comp421/hardware.h`).

4. Project Logistics

4.1. Grading

This project will be graded primarily on correctness and efficiency. However, significant lack of documentation or elegance within your code will adversely affect your grade. You should use reasonable variable names and include comments in particularly tricky parts of the code, but you need not document code that does what it looks like.

4.2. Turning in Your Project

Your project is due by *11:59PM, Tuesday, March 20, 2012*. The project should be done in groups of 2 students, and you should only submit one project for your group (only one group member should submit the project for the group). Everything you want to turn in for your group should be in a single directory (or its subdirectories). Before turning in your project, please create a file named “README” in the directory where your files for the project are located. In this file, *please list the names and NetIDs of both members in your project group*.

Please also describe in your “README” file anything you think it would be helpful for the TAs to know in grading your project. This might, for example, mean describing unusual details of your algorithms or data structures, and/or describing the testing you have done and what parts of the project you think work (or don’t work).

Actually submitting your project for grading will be done similarly to what was done for Lab 1. Specifically, to submit your project for grading when you are ready, one of the group members should perform the following two steps:

- First, on CLEAR, change your current directory to the directory where your files for the project for grading are located. For example, use the “cd” command to change your current directory. When you run the submission program, it will submit everything in (and below) your current directory, including all files and all subdirectories (and everything in them, too). Please make sure all files you want us to see for grading are located in this single directory and/or in subdirectories of this single directory.
- Second, on CLEAR, run the submission program

```
/clear/courses/comp421/pub/bin/lab2submit
```

This program will check with you that you are in the correct directory that you want to submit for grading, and finally, will normally just print “SUCCESS” when your submission is complete. If you get any error messages in running lab2submit, please let me know.

4.3. Honor Code Policy

All assignments in the course are conducted under the Rice Honor Code. For programming assignments such as this one, students are encouraged to talk to each other, the TAs, the instructor, or anyone else about the assignment. This assistance, though, is limited to discussion of the problem; *each project group must produce their own solution*. Consulting another student’s or group’s solution (even from a previous COMP 421 class) is prohibited, and submitted solutions may not be copied from any source.

4.4. Recommended Project Checkpoint

This is a large project, with a significant amount of time to work on it before the due date. We offer the following advice to help you work your way successfully through the project:

- Start working on the project early (like now, if you haven't started already).
- Try to work on the project consistently throughout this time.

If you have any questions, ask a TA or the instructor.

As a further aid in helping you complete the project by the due date, the schedule of important dates for this project includes a recommended *project checkpoint* as an intermediate milestone point within the project. By the date listed above, you should have completed a substantial portion of the project. As a *minimal* guideline, by this date, we suggest that you should have created a Yalnix kernel that can:

- Execute `KernelStart` successfully, initializing the kernel and the machine, including initializing interrupts, enabling virtual memory, and creating the *idle* and *init* processes.
- Be able to execute a simple *init* process that only performs some simple kernel call that blocks, such as the `Delay` kernel call.
- Be able to context switch to the *idle* process while *init* is blocked, and be able to context switch back to *idle* when it becomes unblocked (e.g., when its `Delay` has been completed).

Note that this should be a *minimum* goal to achieve by the checkpoint deadline, with only 2 weeks left in the project. You are *strongly* encouraged to get to this point well before then, and to complete more of the project by the checkpoint date.

The project checkpoint will not be graded, and there is nothing to turn in for it. This schedule and goal for the checkpoint is for your own use in working on the project. You are ultimately responsible for completing the entire project by the overall due date, but if you have not completed the checkpoint on schedule, you will find it very hard to finish with then only 2 weeks remaining. Remember, start on the project early and try to work on it consistently throughout this time.

5. Your Assignment

To review, in order to complete the project, you must implement a Yalnix kernel that runs on the provided RCS 421 computer simulation. Specifically, you must implement:

- a `KernelStart` routine to perform kernel initialization;
- a `SetKernelBrk` routine to change the size of kernel memory;
- a procedure to handle each defined type of interrupt, exception, or trap; and
- a procedure (called by your `TRAP_KERNEL` handler) to implement each defined Yalnix kernel call.

Note that the kernel call prototypes given in Section 3.1 and in `comp421/yalnix.h` are the form in which Yalnix user-level programs use these kernel calls. The function names and prototypes of the procedures that implement them inside your need not be the same. The code in your `TRAP_KERNEL` handler calls whatever functions inside your kernel it wants to, with whatever arguments it wants to, in order to provide the effect of each type of kernel call.

The recommended structure for your kernel would be to make one function in your kernel for each type of kernel call, but to pick some other name for that function inside the kernel. For example, you could name

the function you use to implement the `Delay` inside your kernel `KernelDelay`, and you could then pass any arguments to it and get any return value from it you decide is convenient in your own implementation. Since this procedure is entirely internal to your kernel, called only by your `TRAP_KERNEL` handler, there is no confusion between this and the `Delay` function called by user processes outside the kernel to invoke this function and to generate the trap necessary to enter your kernel to request it.

In addition, the following requirements apply to this project:

- Your solutions must be implemented in C or C++.
- Your solutions must run in the CLEAR/Linux environment, using the supplied RCS 421 computer system simulation environment, as described in Section 2.
- Your kernel must provide the interface and features described in Section 3.
- The project must be done in groups of two students.
- Your kernel must compile and link to an executable program named `yalnix` that is set up to run with the simulated RCS 421 environment provided.
- Your project directory must include a `Makefile` that can be used to make your kernel and any Yalnix user-level test programs you used in testing your kernel. A template `Makefile` is available as `Makefile.template` in the directory `/clear/courses/comp421/pub/templates`. You should copy this file to your project directory and edit it as described in the comments in the file. *In particular, this template Makefile contains special rules for compiling and linking programs, which must be used for the project to work.*

6. Some Implementation Hints

6.1. Process IDs

Each process in the Yalnix system must have a unique integer process ID. Since an unsigned integer allows over 4 billion processes to be created before overflowing its range, you may (for this project) simply assign sequential numbers to each new process created and not worry about integer overflow (real operating system kernels must worry about this, though).

6.2. Fork

On a `Fork`, a new process ID is generated and a new PCB is allocated. The kernel stack (which contains the saved exception stack frame from when the user executed the `TRAP_KERNEL` to call `Fork`) and saved kernel context for the child process are created as copies of the parent (calling) process. New physical memory is allocated for the child process, into which you copy the parent's address space contents. Since the CPU can only access memory by virtual addresses (using the page tables), you must map both the source and the destination of this copy into the virtual address space at the same time. You need not map all of both address spaces at the same time, however: the copy may be done piece-meal, since the address spaces are already naturally divided into pages.

6.3. Exec

On an `Exec`, you must load the new program from the specified (Unix) file into Region 0 of the calling process, which will in general also require changing the process' page tables. The program counter in the `pc` field of the exception stack frame must also be initialized to the new program's entry point. The

template `LoadProgram` function described in Section 3.7 shows how to do this and takes care of most of the messy details for you.

The template `LoadProgram` also shows how to initialize the stack for the new program. The stack for a new program starts with only one page of physical memory allocated, which should be mapped to virtual address `USER_STACK_LIMIT - PAGE_SIZE`. As the process executes, it may require more stack space, which can then be allocated as usual for the normal case of a running program.

To complete initialization of the stack for a new program, the argument list from the `Exec` must first be copied onto the stack. The template `LoadProgram` function also takes care of most of this for you. The stack pointer in the `sp` field of the exception stack frame structure should be initialized by `LoadProgram`. Like all addresses that you use, this must be a virtual address. Again, all these details are presented in the form of the C procedure `LoadProgram` in the file `load.template`.

6.4. Suggested Plan of Attack

You may choose to implement the various parts of your kernel in any order you see fit. If you are having difficulties putting together a coherent plan, we suggest you start with the following sketch of a plan. Please remember to thoroughly test your code using your own stubs or test procedures, and as part of the whole system, at each and every opportunity. The plan of attack below is not necessarily complete, but it is suggestive of a full plan that can allow you to complete the project in some order:

- Read this document carefully. Try to understand the details described in it, and perhaps read this document again. Ask a TA or the instructor if you don't understand parts of this document. Although there are a lot of pages here to read, this document is intended to be very complete and to guide you through the project.
- Think through and sketch out some high-level pseudo-code for each type of kernel call, interrupt, and exception. Then decide on the things you need to put in what data structures (notably in the Process Control Block) to make it all work. Iterate until the pseudo-code and the main prototype data structures mesh well together. Do not worry whether you have everything you will eventually need defined in your PCB structure; you can always add fields to the PCB later (although you must be careful that you recompile everything that uses the PCB definition if you add or change fields in the PCB structure definition).
- Take a cursory look at the template `LoadProgram` function in the file `load.template`. You need not understand all the details, but make sure you understand the comments that are preceded by ">>>>" markers.
- Write an initial version of your `KernelStart` function to bootstrap the kernel. You need to initialize the interrupt vector table here and point the `REG_VECTOR_BASE` register at it. You also need to build the initial page tables and enable virtual memory by writing a 1 into the `REG_VM_ENABLE` register. The support code that makes the RCS 421 simulation work does a lot of error checking at this point, so if you get through this far with no errors, you are probably doing OK (although we can't check for all possible errors at this point).
- Write an *idle* program (a single infinite loop that just executes a `Pause` instruction). Try running your kernel with just an *idle* process to see if that much works.
- Write a simple *init* program. The simplest *init* program would just loop forever. Make sure you can get your `Makefile` to compile this as a Yalnix user program, since you will need this skill in order to write other test programs later. Modify your `KernelStart` to start this *init* program (or one passed in the Unix shell command line) in addition to the *idle* program.

- Make sure your *init* and *idle* processes are context switching. At this point you will be able to pass the recommended checkpoint.
- Implement the `GetPid` kernel call and call it from the *init* process. At this point your kernel call interface would seem to be working correctly.
- Implement `SetKernelBrk` to allow your kernel to allocate substantial chunks of memory. It is likely that you haven't needed it up to this point, but you may have (in this case implement it earlier).
- Implement the `Brk` kernel call and call it from the *init* process. At this point you have a substantial part of the memory management code working.
- Implement the `Delay` kernel call and call it from the *init* process. Make sure your *idle* process then runs for several clock ticks uninterrupted, until the delay period expires. This will be the first proof that blocking of processes works.
- Implement the `Fork` kernel call. If you get this to work you are almost done with the memory system.
- Implement the `Exec` kernel call. You have already done something similar by initially loading *init*.
- Write another small program that does not do much. Call `Fork` and `Exec` from your *init* process, to get this third program running as a child of *init*. Watch for context switches.
- Implement and test the `Exit` and `Wait` kernel calls.
- Implement the kernel calls related to terminal I/O handling. These should be easy at this point, if you pay attention to the process address space into which your input needs to go.
- Look at your work and wonder in amazement at the uncertain, frustrating, and rewarding road you have traveled.

7. Compiling, Running, and Using Yalnx

7.1. Compiling Your Kernel

The file `/clear/courses/comp421/pub/templates/Makefile.template` contains a basis for the `Makefile` we recommend you use for this project. The comments within it should be self-explanatory. There is some magic involved how this `Makefile` sets up compiling and linking with `gcc` and `g++`, but you may safely ignore most of the details and just switch between the two compilers with the `"LANG ="` definition at the top of the file. The arguments passed to the compiler and linker to make the kernel and Yalnx user programs should not be changed.

7.2. Running Your Kernel

Your kernel will be compiled and linked as an ordinary Unix executable program, and can thus be run as a command from the Unix shell prompt. When you run your kernel, you can put a number of Unix-style switches on the command line to control some aspects of execution. The file name for your executable Yalnx kernel should be `yalnx`. You can then run your kernel as:

```
yalnx -t tracefile -lk level -lh level -lu level -s -n initfile initargs...
```

All of the switches and arguments shown here on the command line are optional. The meaning of each of the command-line switches is as follows:

- `-t tracefile`: This turns on “tracing” within the kernel and machine simulation code, and optionally specifies the name of the file to which the tracing output should be written. If no *tracefile* argument follows the `-t` switch, the default *tracefile* name is “TRACE” in the current directory.

To generate trace output from your kernel, you may call

```
TracePrintf(int level, char *fmt, args...)
```

where `level` is an integer tracing level, `fmt` is a format specification in the style of `printf`, and `args...` are the arguments needed by `fmt`. You can run your kernel with the “tracing” level set to any integer. If the current tracing level is greater than or equal to the `level` argument to `TracePrintf`, the output generated by `fmt` and `args...` will be added to the trace. Otherwise, no trace output is generated from this call to `TracePrintf`.

- `-lk level`: Set the tracing level for the kernel for this run. The default tracing level is -1, if you enable tracing with the `-t` or `-s` switches. You can specify any level of tracing.
- `-lh level`: Like `-lk`, but sets the trace level to be applied to internal `TracePrintf` calls made by the hardware simulation. The higher the number, the more verbose, complete, and incomprehensible the hardware trace. Tracing the hardware simulation may sometimes be useful to you in your own debugging; if you encounter any really strange problems, this tracing may be very useful to us in helping you find your problem or in diagnosing any internal simulator problems (although no such problems are expected).
- `-lu level`: Like `-lk` and `-lh`, but sets the tracing level applied to `TracePrintf` calls made by user-level processes running on top of Yalnx.
- `-n`: Do not use the X window system support for the simulated terminals attached to the Yalnx system. The default is to use X windows.
- `-s`: Send the tracing output to the Unix `stderr` file (this is usually your screen) in addition to sending it to the tracefile. This switch enables tracing, even if the `-t` switch is not specified.

These switches are automatically parsed, interpreted, and deleted for you before `KernelStart` is called. The remaining arguments from the command line are passed to `KernelStart` in its `cmd_args` argument. For example, if you run your kernel with the command line

```
yalnx -lk 5 -lh 3 -n init a b c
```

then your `KernelStart` would be called with `cmd_args` as follows:

```
cmd_args[0] = "init"
cmd_args[1] = "a"
cmd_args[2] = "b"
cmd_args[3] = "c"
cmd_args[4] = NULL
```

Inside `KernelStart`, you should use `cmd_args[0]` as the file name of the *init* program, and you should pass the whole `cmd_args` array as the arguments for the new process. You should use a default *init* program name of “init” if no `cmd_args` are specified when Yalnx is run from the Unix shell.

7.3. Debugging Your Kernel

You may use `fprintf`, `printf`, and `TracePrintf` within your kernel to generate any necessary debugging output you want to see; we recommend that you use `TracePrintf`. Please be aware that the insertion of such print statements may alter the behavior of your program, since they can change the timing of when things happen relative to when clock interrupts or terminal interrupts happen, for example.

You may also use `gdb` to debug your kernel. However, because of the RCS 421 simulator's heavy use of Unix signals internally, it is necessary to tell `gdb` to ignore certain events:

- For `gdb`, you can do this by typing the following at the `gdb` prompt:

```
handle SIGILL SIGFPE SIGBUS SIGSEGV pass nostop noprint
handle SIGALRM SIGUSR1 SIGCHLD SIGPOLL pass nostop noprint
```

Putting these commands in your `.gdbinit` file will relieve you from needing to type them every time you run `gdb` to debug your kernel.

7.4. Controlling Your Yalnix Terminals

By default (unless you specify the option `-n` option on the command line or unless you are not running X windows) each Yalnix terminal is simulated as an `xterm` on your X windows display.

The current RCS 421 machine configuration supports four terminals, numbered from 0 to 3, with terminal 0 serving as the Yalnix system console, and the other 3 serving as regular terminals. In fact, though, these uses are only a convention, and all four terminals actually behave in exactly the same way. The constant `NUM_TERMINALS` in `comp421/hardware.h` defines the total number of terminals supported. The constant `TERMINAL_MAX_LINE` defines the maximum length of line supported for either input or output on the terminals.

When you run your kernel, the X window system will create four new windows on your display, called

```
Yalnix Console
Yalnix Terminal #1
Yalnix Terminal #2
Yalnix Terminal #3
```

to act as the four terminals that the hardware configuration supports. Each of these windows is actually running a Unix `xterm` and can be controlled as such. For example, you can resize or iconify any of these windows, and can use the scroll bar to review output in the window. For those who like to play with configuring X windows, you can define the shell environment variable `YALNIX_XTERM` to contain `xterm` command line arguments to be used when starting the Yalnix terminal windows. You can also define the shell environment variables `YALNIX_XTERM0`, `YALNIX_XTERM1`, `YALNIX_XTERM2`, and `YALNIX_XTERM3` to be further command line arguments to be used only when starting the corresponding terminal window. For example, you could define these four environment variables to contain appropriate “-geometry” options to automatically place each window on the screen for you. You can also include `xterm` options to change the font size in each window, change the foreground or background colors, etc.

When your kernel program ends, these four X windows automatically go away. However, to allow you to look at them before they disappear, the `Halt` instruction pauses before exiting from your kernel program, and prints

```
Press RETURN to end simulation...
```


Once you hit the RETURN key, the machine simulation will exit and the console and terminal windows will disappear.

The terminal windows keep log files of all input and output operations on each terminal. The files `TTYLOG.0`, `TTYLOG.1`, `TTYLOG.2`, and `TTYLOG.3` record all input and output from each terminal separately; and the file `TTYLOG` collectively records all input and output from the four terminals together. The terminal logs show the terminal number of the terminal, either “<” for input or “>” for output, and the line input or output. You can type an end of file on a terminal with control-D as in Unix, which appears in the terminal log files as “(EOF)”.

Table of Contents

1	Project Overview	1
2	The RCS 421 Computer System	1
2.1	Machine Registers	2
2.1.1	General Purpose Registers	2
2.1.2	Privileged Registers	2
2.2	Memory Subsystem	3
2.2.1	Overview	3
2.2.2	Physical Memory	4
2.2.3	Virtual Address Space	4
2.2.4	Page Tables	5
2.2.5	Translation Lookaside Buffer (TLB)	7
2.2.6	Enabling Virtual Memory	8
2.3	Hardware Devices	8
2.3.1	Terminals	8
2.3.2	The Hardware Clock	9
2.4	Interrupts, Exceptions, and Traps	10
2.4.1	Interrupt Vector Table and Types	10
2.4.2	Hardware-Defined Exception Stack Frame Structure	11
2.5	Additional CPU Machine Instructions	12
3	Yalnix Operating System Details	12
3.1	Yalnix Kernel Calls	12
3.2	Interrupt, Exception, and Trap Handling	15
3.3	Kernel Booting Entry Point	17
3.4	Memory Management	17
3.4.1	Virtual Address Space Layout	18
3.4.2	Kernel Memory Management	19
3.4.3	Initializing the Page Tables and Virtual Memory	20
3.4.4	Growing a User Process's Stack	22
3.5	Context Switching	22
3.6	Bootstrapping and Kernel Initialization	24
3.7	Loading Yalnix User Programs from Disk	25
3.8	Additional Terminal Interface	25
4	Project Logistics	26
4.1	Grading	26
4.2	Turning in Your Project	26
4.3	Honor Code Policy	26
4.4	Recommended Project Checkpoint	27
5	Your Assignment	27

6	Some Implementation Hints	28
6.1	Process IDs	28
6.2	Fork	28
6.3	Exec	28
6.4	Suggested Plan of Attack	29
7	Compiling, Running, and Using Yalnx	30
7.1	Compiling Your Kernel	30
7.2	Running Your Kernel	30
7.3	Debugging Your Kernel	32
7.4	Controlling Your Yalnx Terminals	32